

# Identifying blocks/records on disks and in memory

- Identifying a block/record on *disk*: Block Address and record address

- Fact:

- A **block** is **identified** by a **block address**
  - A **record** is **identified** by a **record address**

(A **record address** is an **extension** of a **block address**, so in this discussion, I will often use: **block/record address** and handle **both topics** at the same time)

- Identifying a block/record in **memory**

- Fact:

- A **block/record** can be **read** into the **computer memory**

- Identifying a **block/record** when it is **in memory**:

- Use a **virtual memory address** (**paging**)


- Overview: Types of addresses to identify blocks/records**

- There are **2 types** of **address** to **identify** a **block/record** in use:

- Database Address**: used to identifies **data (block or record)** stored on **disk**  
  
There are **2 kinds** of **database addresses**:

- Physical address**
    - Logical address**
  - Virtual memory address**: used to identify **data (block or record)** stored in **(virtual) memory**

Graphically:



- Database address = address used to identify a block (data) on disk**

- Database Address:

- Database Address** = an **address** used to identify **data (block/record)** stored on **disk**

- The are **2 common techniques** to **identify** a **block/record** stored on **disk**:

- **Physical address**
  - **Logical address**

- **Physical Address**: a *direct* addressing format for identify block/record on a disk

- Physical Address of a **block**:

- **Physical Address** of a **block** = **address** used to identify **block** on a **disk**

---

- **Components** of a **block physical address**:

disk 1

disk 2

disk 1

disk 2

Host 1

Host 2

- **Host ID** that contains the **disk**
- **Disk ID** the **identifies** the disk on the **host**
- **Block (cylinder/track/section) number** on the **disk**

- Physical Address of a **record**:

- **Physical Address** of a **record** = **address** used to identify **record** on a **disk**

---

- **Components** of a **record physical address**:

- **Physical address** of a **block**
- **Offset** of the **record** in the **block**

- **Logical Address**: an *indirect* addressing format for identify block/record on a disk

- Logical block address:

- Each **block/address** is assigned a unique **logical address**

---

- **Logical address** = an **arbitrary string** of fixed length **bits**

(Can be generated **automatically** using some sequence generator or keep adding 1 to a counter)

- DBMS uses a **map table** to **translate**:

- **Map table**:

Logical Address	Physical Address
101001.....10	(HostID, DiskID, ..., BlockID)
010101.....01	(HostID, DiskID, ..., BlockID)
...	

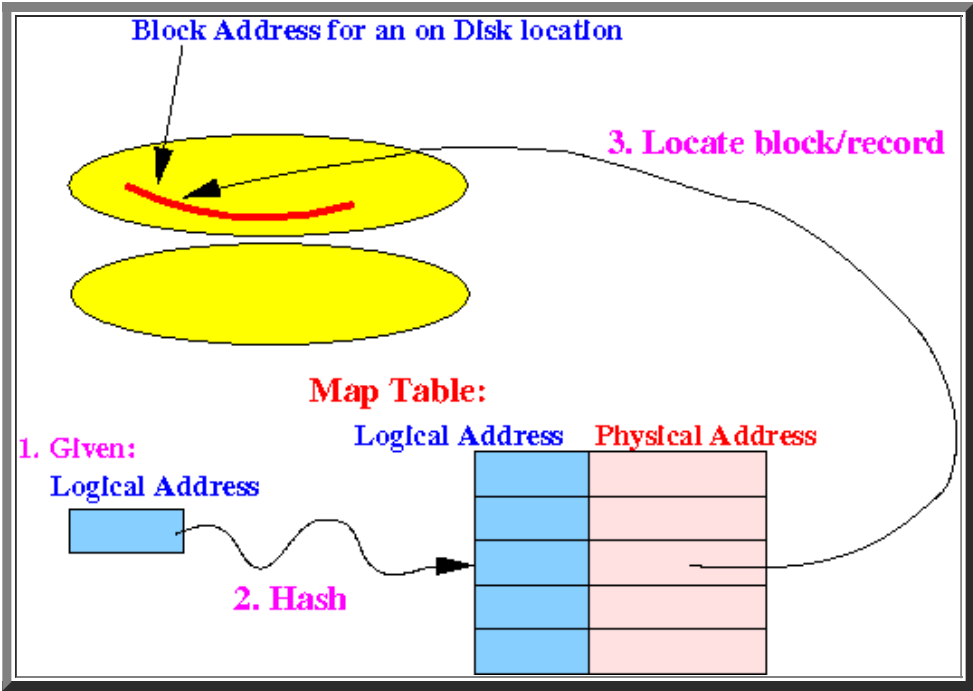
- To **speed up** access, the **Map Table** is organized as:

- a **hash table**.

---

- The **map table** is stored on **disk** (in a well-known location)

- **How** to use the **map table**:



- *Usage of (physical/logical) database addresses*
  - **How** a **(physical/logical) database address** is **used**:

- When some **record *x*** need to **refer** to *another record *y**, we can **store**:
 

- the **(physical/logical) database address** of *y* inside *x*

- *Motivation for using logical addressing*
  - **Intro**: Techniques to **reference (point)** to a **record**:


- **Primary key/foreign key**:
 

- Each **record *x*** contains a **unique primary key**
  - Another record *y* that **references** the **record *x*** will store the **primary key** of *x*  
(The **primary key** of *x* stored in **record *y*** is called a **foreign key**)
- The **direct approach**:
 

- The **record *y*** stores the **physical/logical address** of the **record *x***

- **Problem** with **referencing** another **record** using a **physical address**:

- **Example**: 2 records references the **record *Y***

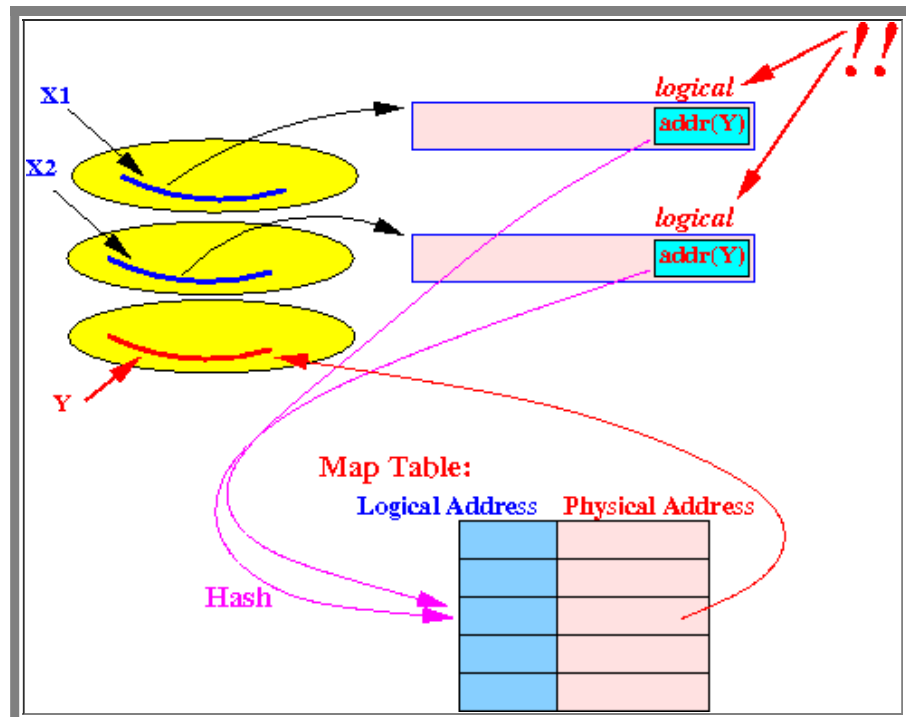


- **Problem:** If **record Y** is **moved** to a **different part** of the **disk**:

- We must **update many** addresses !!!

- **Solution** using a **logical address** as **reference**:

- **Example:** 2 records references the **record Y**



- If we **move** the **record Y**, we **only** need to:

- **Update** the **physical address** in the **map table** !!!

- **Caveat: block in memory**

- **Fact:**

- A **block/record** is stored on **disk** (this is **not** the **problem**)

- **Caveat/Problem:**

- A **block/record** must be stored in **main memory** in order to perform **read/write operations** on **data** in the **block**

- **Virtual Address:** address used to identify block/record in main memory

- Main memory:

■ Each data item stored in (main) memory is identified by its **virtual memory address**

■ While a **block/record** is in **memory**, the DBMS must use the **virtual address** to identify the **block/record**:

Main memory:

Virtual address of block in memory

Virtual address of record in memory

Block

Record

- Common ways to identify a record inside a block

- Consider the **records** stored in a **block**:

Store offsets in block header

Offsets of each record in block

Block X: 

block header

o1

o2

record 2

record 1

0 1

index

Clearly:

■ Address of a **record** = address of the **block** that contains the **record** + some **offset** information

- There are **2 common ways** to **identify** a record:

1. (DatabaseAddress(Block X) + the **(absolute) offset in block**)

Example:

■ record 1: (DatabaseAddress(Block X), 234 (offset in bytes))

■ record 2: (DatabaseAddress(Block X), 456 (offset in bytes))

2. (DatabaseAddress(Block X) + the **index** of the **offset table** (in the block header))


Example:

■ record 1: (DatabaseAddress(Block X), 0)


■ record 2: (DatabaseAddress(Block X), 1)

- A caveat when using **physical** addresses to reference a block/record

- **Illustrative example:** 2 block with **records** referencing a **record Y**:



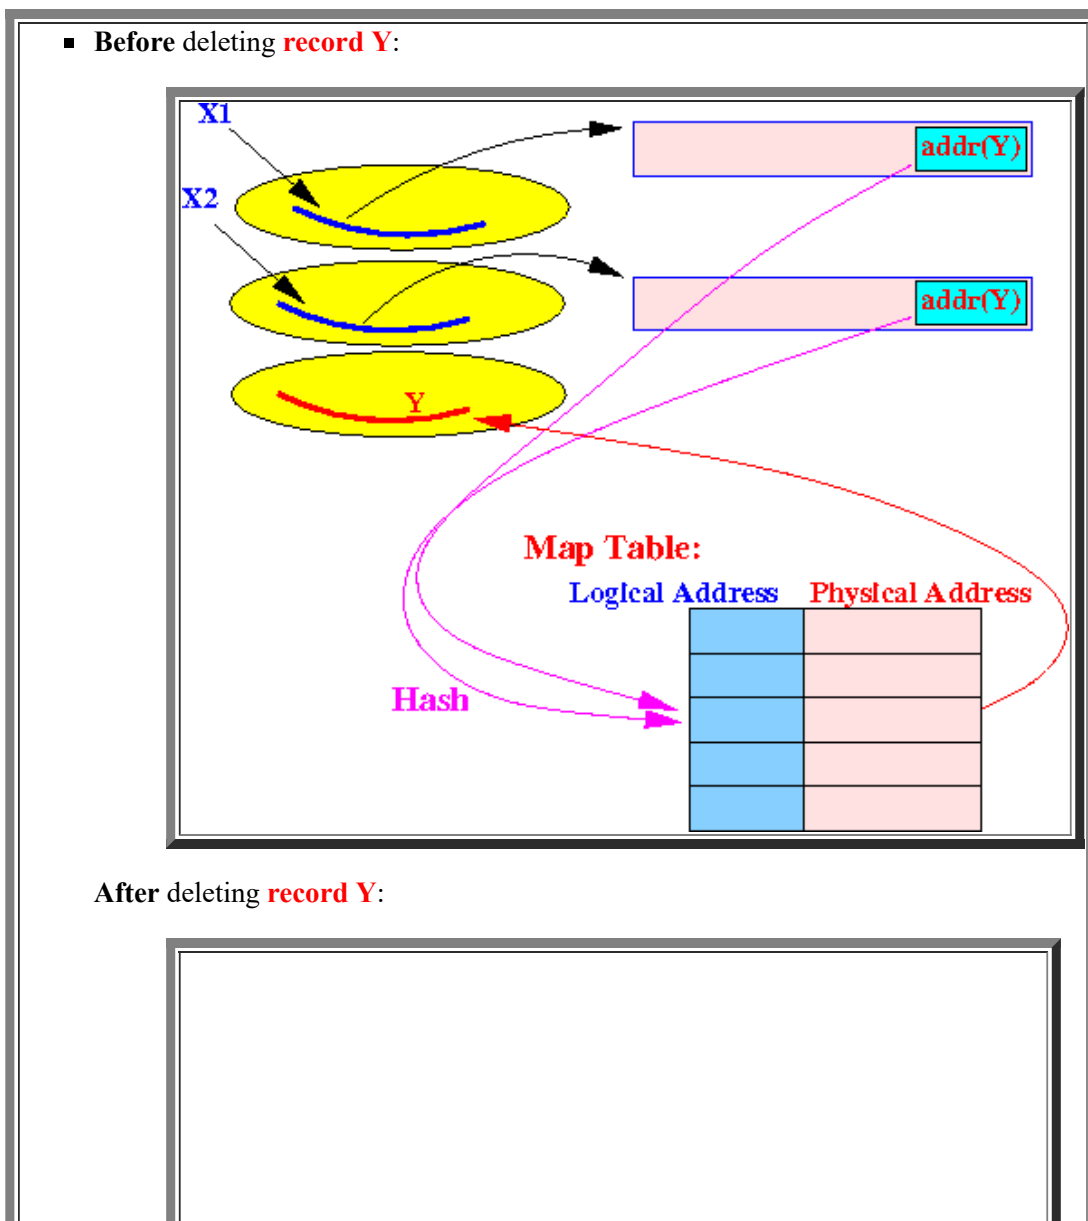
- When the **record Y** is **deleted** :

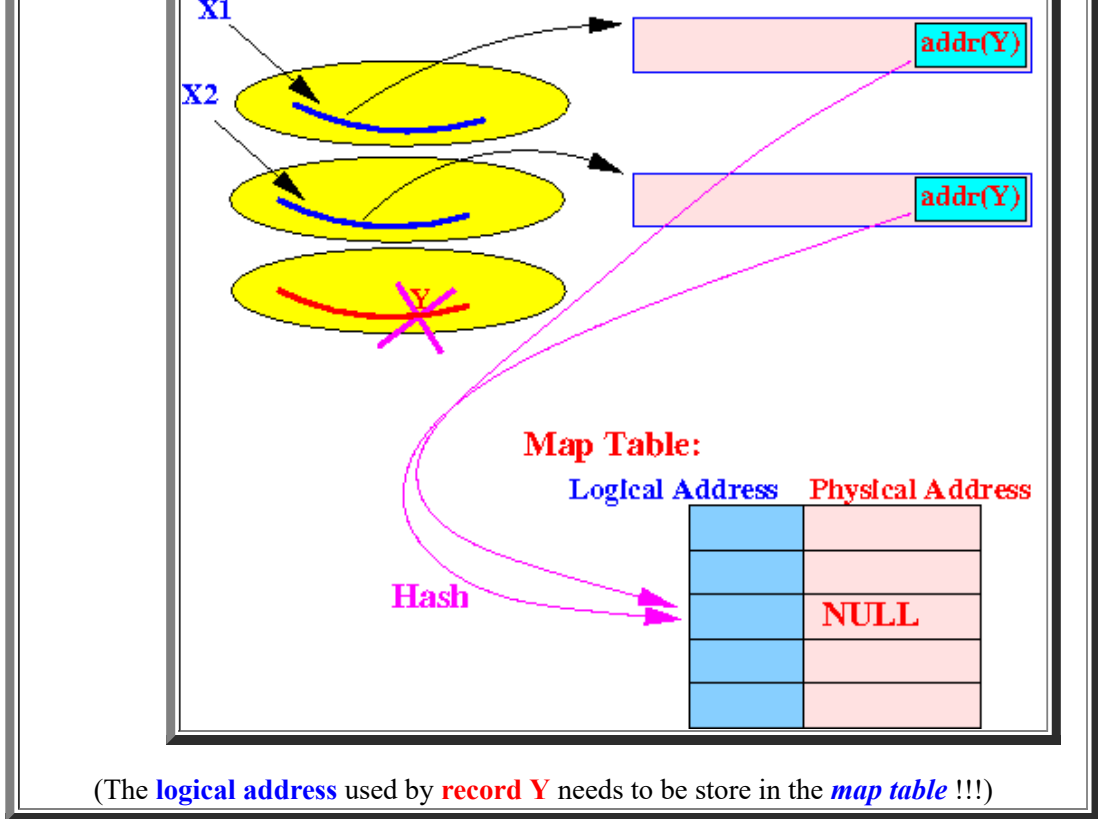


the *physical* addresses will reference an **incorrect record** !!!

### Techniques to handle record deletion

- Using *logical* addresses is **easy**:

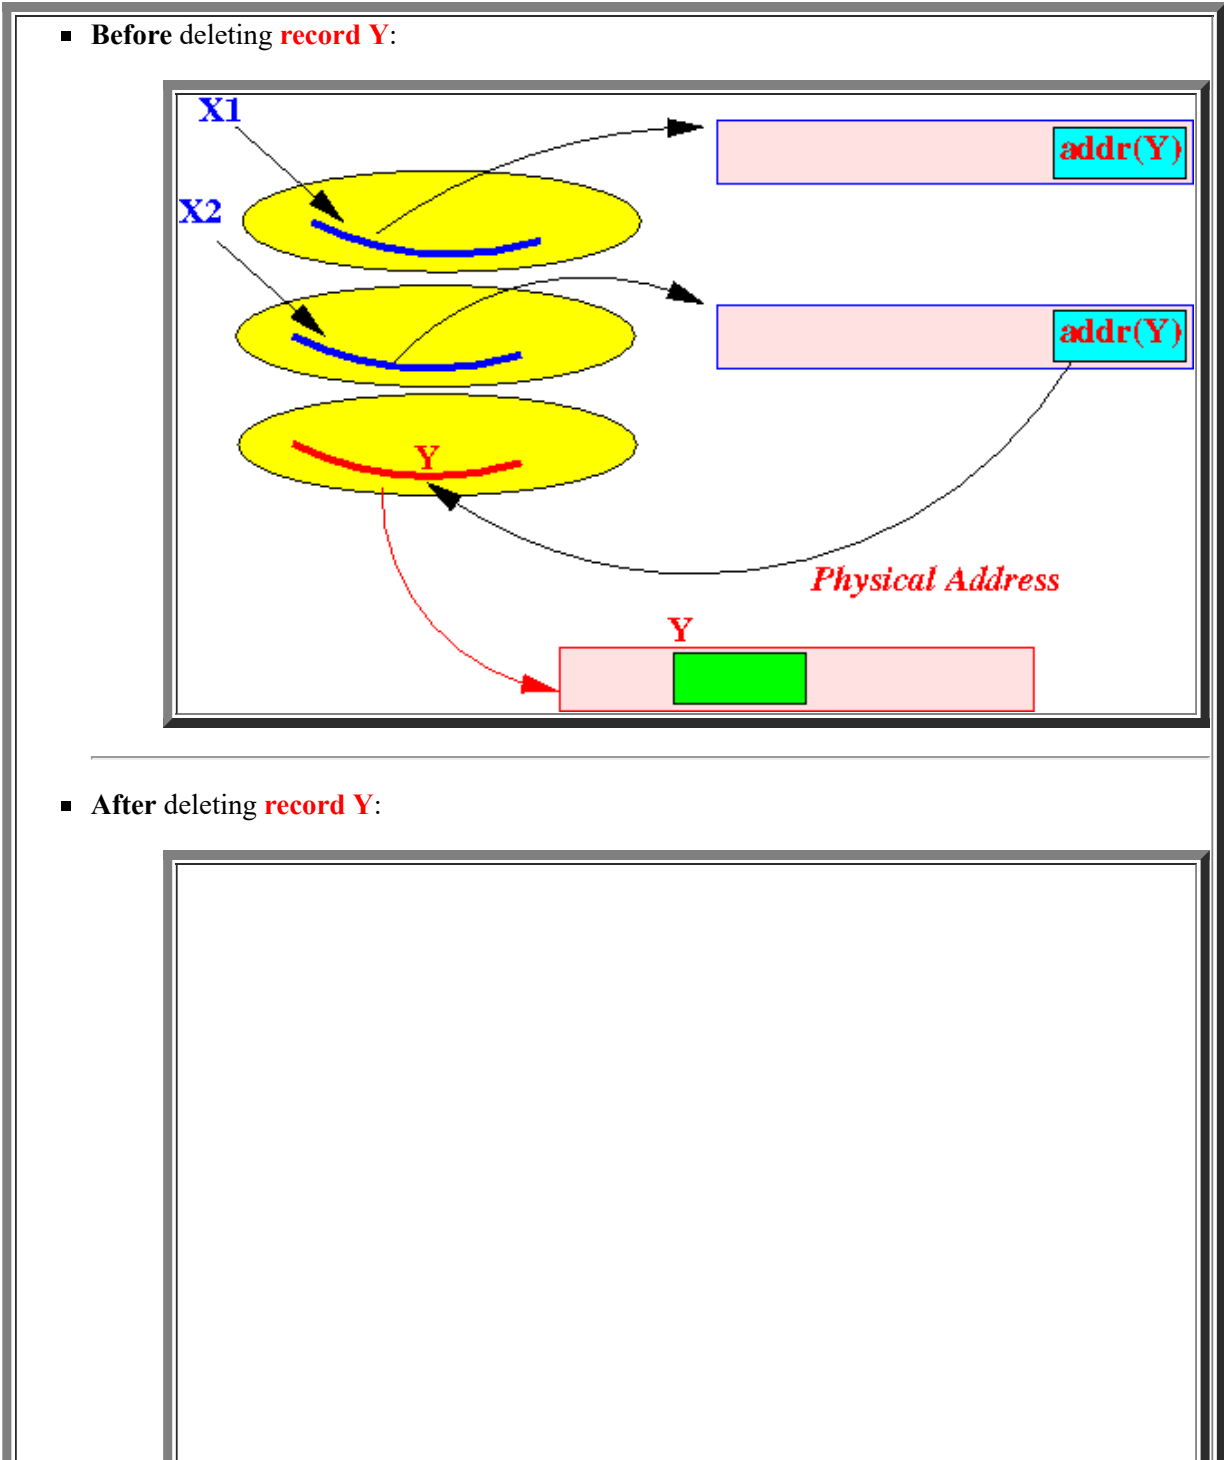


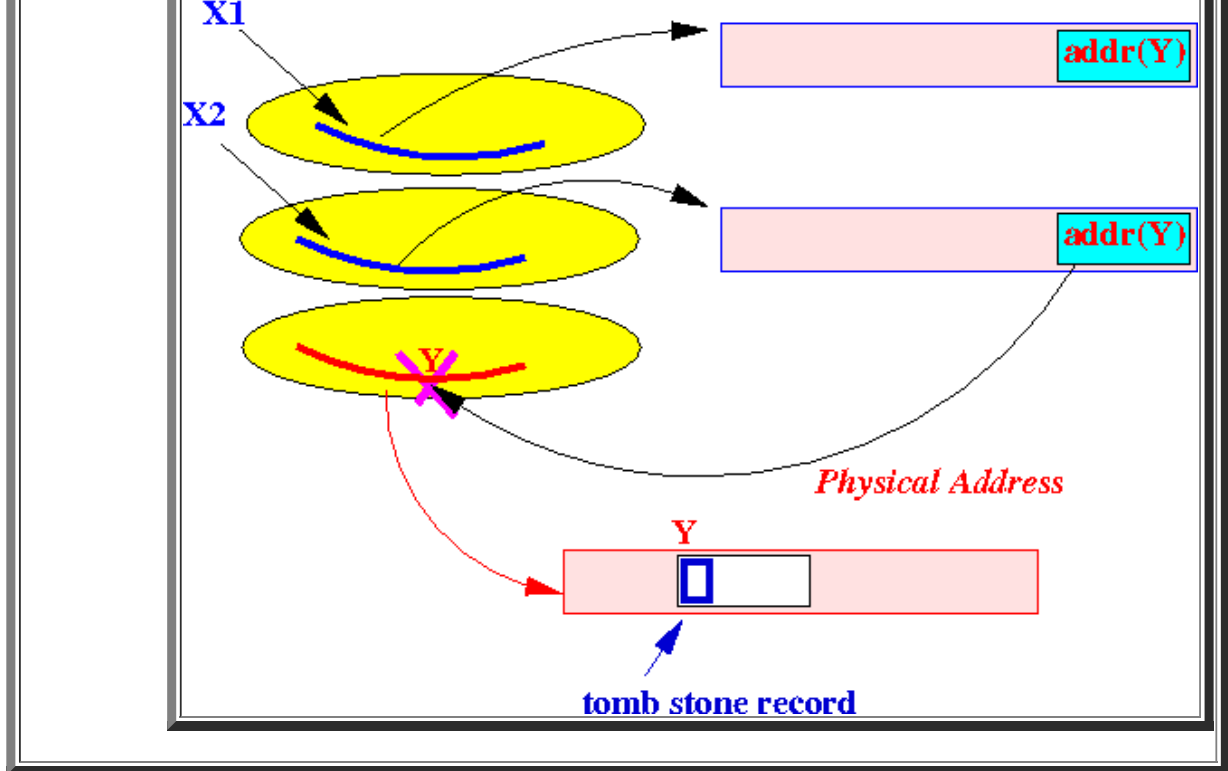


- Deleting a **record** using **physical address**: use a **tomb stone record**

- **Tomb stone record**: a (very small) **special purpose record** used to indicate a **deleted record**
- When a **record** is delete, it is **replaved** by the **tomb stone record**

Example:





• Caveat using a tomb stone record

◦ Insertion:

- When you **insert** a new **record**:

■ You **cannot** use the **space** of a **tomb stone record** !!!

Because:

- **Existing** record references to the **deleted record** will **then** references to the **newly inserted record**:

